

Optimization

Pieter Vanderschueren

Year 2024-2025

Contents

Mathematical Preliminaries	2
1 Vectors	3
2 Independence, Subspaces, Basis and Dimension	4
3 Orthogonality and Orthogonal Complements	5
4 Matrices	6
5 Sequences	10
6 Differential Calculus	11
Fundamental Concepts	14
7 Fundamental Concepts of Optimization	15
8 Types of Optimization Problems	16
9 Convex Optimization	19
10 The Lagrangian Function and Duality	25
Unconstrained Optimization Algorithms	30
11 Optimality Conditions	31
12 Estimation and Fitting Problems	33
13 Newton Type Optimization	34
14 Globalisation Strategies	34
15 Calculating Derivatives	34
Constrained Optimization Algorithms	35
16 Optimality Conditions for Constrained Optimization	36
17 Equality Constrained Optimization Algorithms	36
18 Inequality Constrained Optimization Algorithms	36
19 Optimal Control Problems	36

Mathematical Preliminaries

1 Vectors

Theorem 1.1: Vector Space \mathbb{R}^n

The vector space \mathbb{R}^n is the set of all n -dimensional column vectors with real components. The space \mathbb{R}^n is equipped with the following operations:

- component-wise addition:

$$x + y = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} + \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} x_1 + y_1 \\ x_2 + y_2 \\ \vdots \\ x_n + y_n \end{bmatrix}$$

- scalar multiplication:

$$\alpha x = \begin{bmatrix} \alpha x_1 \\ \alpha x_2 \\ \vdots \\ \alpha x_n \end{bmatrix}$$

Theorem 1.2: Dot product

The dot product of two vectors $x, y \in \mathbb{R}^n$ is defined as the scalar

$$x \cdot y = x^T y = \sum_{i=1}^n x_i y_i.$$

Theorem 1.3: Norm

A norm $\|\cdot\|$ on \mathbb{R}^n is a function $\|\cdot\| : \mathbb{R}^n \rightarrow \mathbb{R}$ satisfying the following:

- non-negativity: $\|x\| \geq 0$ for any $x \in \mathbb{R}^n$; $\|x\| = 0$ if and only if $x = 0$,
- positive homogeneity: $\|\alpha x\| = |\alpha| \|x\|$ for any $x \in \mathbb{R}^n$ and $\alpha \in \mathbb{R}$,
- triangle inequality: $\|x + y\| \leq \|x\| + \|y\|$ for any $x, y \in \mathbb{R}^n$.

Theorem 1.4: ℓ_p -norms

The class of ℓ_p -norms is defined by

$$\|x\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{\frac{1}{p}}$$

which includes the following special cases:

- ℓ_1 -norm: $\|x\|_1 = \sum_{i=1}^n |x_i|$,
- ℓ_2 -norm: $\|x\| = \|x\|_2 = \left(\sum_{i=1}^n x_i^2 \right)^{\frac{1}{2}}$,
- ℓ_∞ -norm: $\|x\|_\infty = \max_{i=1, \dots, n} |x_i|$.

Theorem 1.5: Angle between two vectors

The angle $\angle(x, y)$ between two vectors $x, y \in \mathbb{R}^n$ is given by

$$\angle(x, y) = \arccos\left(\frac{x \cdot y}{\|x\|\|y\|}\right).$$

We say that two vectors $x, y \in \mathbb{R}^n$ are

- orthogonal if $\angle(x, y) = \frac{\pi}{2}$, i.e. $x \cdot y = 0$,
- aligned if $\angle(x, y) = 0$,
- anti-aligned if $\angle(x, y) = \pi$,
- parallel if $\angle(x, y) = 0$ or $\angle(x, y) = \pi$,
- at an acute angle if $\angle(x, y) < \frac{\pi}{2}$,
- at an obtuse angle if $\angle(x, y) > \frac{\pi}{2}$.

Theorem 1.6: Cauchy-Schwarz inequality

For any two vectors $x, y \in \mathbb{R}^n$, the following inequality holds:

$$|x \cdot y| \leq \|x\|\|y\|.$$

Theorem 1.7: Hölder inequality

Let $p \geq 1$. For any $x, y \in \mathbb{R}^n$, the following inequality holds:

$$|x \cdot y| \leq \|x\|_p \|y\|_q,$$

where $q = \frac{p}{p-1}$ is the Hölder conjugate of p .

2 Independence, Subspaces, Basis and Dimension

Theorem 2.1: Linear Independence

A set of vectors a_1, a_2, \dots, a_n in \mathbb{R}^n is said to be linearly independent if no vector in the collection can be expressed as a combination of the others. In other words

$$\sum_{i=1}^m \lambda_i a_i = 0 \Rightarrow \forall i \in [1, m] : \lambda_i = 0$$

Theorem 2.2: Subspace

A nonempty subset S of \mathbb{R}^n is called a subspace if for any real numbers λ_1, λ_2

$$a_1, a_2 \in S \Rightarrow \lambda_1 a_1 + \lambda_2 a_2 \in S$$

A subspace always contains the zero element.

Theorem 2.3: Span

Given a set of vectors a_1, a_2, \dots, a_n in \mathbb{R}^n , the set of linear combinations of these vectors is called the span, i.e.

$$\text{span}\{a_1, a_2, \dots, a_n\} = \left\{ y \in \mathbb{R}^n \mid y = \sum_{i=1}^n \lambda_i a_i \right\}$$

Theorem 2.4: Basis

The subset $B = \{a_1, a_2, \dots, a_n\}$ of \mathbb{R}^n is called a basis if it is linearly independent and spans \mathbb{R}^n , i.e. it is a maximally independent subset of \mathbb{R}^n .

3 Orthogonality and Orthogonal Complements

Theorem 3.1: Orthogonality

A set of nonzero vectors a_1, a_2, \dots, a_n in \mathbb{R}^n is said to be orthogonal if the dot product of any two distinct vectors is zero, i.e.

$$\forall i, j \in [1, n] : i \neq j \Rightarrow a_i \cdot a_j = 0$$

Consequently, two subspaces S_1 and S_2 of \mathbb{R}^n are orthogonal if every vector in S_1 is orthogonal to every vector in S_2 . In that case, S_2 is called the orthogonal complement of S_1 and denoted by S_1^\perp . The following now holds true for a subspace S of \mathbb{R}^n :

$$\mathbb{R}^n = S \oplus S^\perp$$

Theorem 3.2: Orthonormality

A set of nonzero vectors a_1, a_2, \dots, a_n in \mathbb{R}^n is said to be orthonormal if it is orthogonal and each vector has a unit length, i.e.

$$\forall i, j \in [1, n] : i \neq j \Rightarrow a_i \cdot a_j = 0 \wedge \|a_i\| = 1$$

4 Matrices

Example 4.1: Types of matrices

A matrix $A \in \mathbb{R}^{m \times n}$ is said to be

- the zero matrix, denoted by 0 , if all its entries are zero,
- a square matrix if $m = n$,
- the identity matrix if it is square and all its diagonal entries are one,
- a diagonal matrix if all its off-diagonal entries are zero,
- an upper triangular matrix if all its entries below the diagonal are zero,
- a lower triangular matrix if all its entries above the diagonal are zero,
- a symmetric matrix if it is square and $A = A^T$, the set of these matrices is denoted by \mathbb{S}^n ,
- an orthogonal matrix if it is square and $AA^T = A^T A = I$,
- a non-singular matrix if it is square and there exists another square matrix $B \in \mathbb{R}^{n \times n}$, the inverse of A , such that

$$AB = BA = I$$

- a dyadic matrix if it is of the form $A = uv^T$ for some vectors $u, v \in \mathbb{R}^n$.

Theorem 4.1: Range

Given a matrix $A \in \mathbb{R}^{m \times n}$, the range of A is the set of m -dimensional vectors that can be expressed as Ax for some n -dimensional vector x , and we denote it by

$$\mathcal{R}(A) = \{Ax \mid x \in \mathbb{R}^n\}$$

In other words, it is the set of vectors that can be expressed as linear combinations of the columns of A .

Theorem 4.2: Kernel

Given a matrix $A \in \mathbb{R}^{m \times n}$, the kernel of A is the set of n -dimensional vectors that are mapped to the zero vector by A , and we denote it by

$$\mathcal{N}(A) = \{x \in \mathbb{R}^n \mid Ax = 0\}$$

In other words, it is the set of vectors that are orthogonal to the columns of A .

Theorem 4.3: Rank

Given a matrix $A \in \mathbb{R}^{m \times n}$, the rank of A is the dimension of its range, i.e.

$$\text{rank}(A) = \dim(\mathcal{R}(A))$$

Note: The sum of the dimensions of the range of A and the null space (kernel) of A is equal to the number of columns n :

$$\dim(\mathcal{R}(A)) + \dim(\mathcal{N}(A)) = n$$

Theorem 4.4: Fundamental theorem of Linear Algebra

For any matrix $A \in \mathbb{R}^{m \times n}$, it holds that $\mathcal{N}(A) \perp \mathcal{R}(A^T)$ and $\mathcal{N}(A^T) \perp \mathcal{R}(A)$, therefore we have

$$\mathbb{R}^n = \mathcal{N}(A) \oplus \mathcal{R}(A^T)$$

$$\mathbb{R}^m = \mathcal{N}(A^T) \oplus \mathcal{R}(A)$$

Theorem 4.5: Singular Value Decomposition

Every matrix $A \in \mathbb{R}^{m \times n}$ of rank r can be written as

$$A = U\Sigma V^T = \begin{bmatrix} U_1 & U_2 \end{bmatrix} \begin{bmatrix} \Sigma_1 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} V_1^T \\ V_2^T \end{bmatrix}$$

where $U \in \mathbb{R}^{m \times m}$ and $V \in \mathbb{R}^{n \times n}$ are orthogonal matrices, $U_1 \in \mathbb{R}^{m \times r}$, $V_1 \in \mathbb{R}^{n \times r}$ and

$$\Sigma_1 = \text{diag}(\sigma_1, \sigma_2, \dots, \sigma_r) \in \mathbb{R}^{r \times r}$$

is a diagonal matrix, where $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r > 0$ are the singular values of A . The columns of U and V are called the left and right singular vectors of A , respectively.

Property 4.1: Singular Value Decomposition

The singular value decomposition of a matrix $A \in \mathbb{R}^{m \times n}$ gives orthogonal bases for the four fundamental subspaces related to A :

$$\begin{aligned} \mathcal{R}(A) &= \mathcal{R}(U_1), & \mathcal{N}(A^T) &= \mathcal{R}(U_2) \\ \mathcal{R}(A^T) &= \mathcal{R}(V_1), & \mathcal{N}(A) &= \mathcal{R}(V_2) \end{aligned}$$

Theorem 4.6: Moore Penrose Pseudo Inverse

Assume $J \in \mathbb{R}^{m \times n}$ and that the singular value decomposition (SVD) of J is given by $J = U\Sigma V^T$. Then, the Moore-Penrose pseudo-inverse J^+ is given by

$$J^+ = VS^+U^T,$$

where for

$$S = \begin{bmatrix} \sigma_1 & & & & & & \\ & \sigma_2 & & & & & \\ & & \ddots & & & & \\ & & & \sigma_r & & & \\ & & & & 0 & & \\ & & & & & \ddots & \\ & & & & & & 0 \\ \hline 0 & \dots & \dots & 0 & \dots & \dots & 0 \end{bmatrix} \quad \text{holds} \quad S^+ = \begin{bmatrix} \sigma_1^{-1} & & & & & & 0 \\ & \sigma_2^{-1} & & & & & \vdots \\ & & \ddots & & & & \vdots \\ & & & \sigma_r^{-1} & & & 0 \\ & & & & 0 & & \vdots \\ & & & & & \ddots & \vdots \\ & & & & & & 0 \end{bmatrix}$$

Property 4.2: Moore Penrose Pseudo Inverse

If matrix $A \in \mathbb{R}^{m \times n}$

- is non-singular then

$$A^+ = A^{-1},$$

- has full column rank, that is $r = n$, then

$$A^+A = VV^T = I,$$

i.e. A^+ is a left inverse of A ,

- has full row rank, that is $r = m$, then

$$AA^+ = UU^T = I,$$

i.e. A^+ is a right inverse of A .

Theorem 4.7: Orthogonal-triangular decomposition (QR)

If $A \in \mathbb{R}^{m \times n}$, then there exists an orthogonal matrix $Q \in \mathbb{R}^{m \times m}$ and an upper triangular matrix $R \in \mathbb{R}^{m \times n}$ such that

$$A = QR.$$

Theorem 4.8: Eigenvalue decomposition

Any real symmetric matrix $A \in \mathbb{R}^{n \times n}$ can be decomposed as

$$A = Q\Lambda Q^T,$$

where $Q \in \mathbb{R}^{n \times n}$ is orthogonal, i.e. $Q^T Q = I$, and $\Lambda \in \mathbb{R}^{n \times n}$ is diagonal with the eigenvalues of A on the diagonal. The columns of Q form an orthonormal set of eigenvectors.

Theorem 4.9: Symmetric positive semi-definite matrices

We write for a symmetric matrix $B = B^T$, $B \in \mathbb{R}^{n \times n}$ that “ $B \succeq 0$ ” if and only if B is positive semi-definite, i.e.,

$$\forall z \in \mathbb{R}^n : z^T B z \geq 0,$$

or, equivalently, if all (real) eigenvalues of B are non-negative. The set of all symmetric positive semi-definite matrices is denoted by \mathbb{S}_+^n .

Property 4.3: Symmetric positive semi-definite matrices

Let $Q \in \mathbb{S}^n$ be a symmetric matrix. Then the following statements are equivalent:

1. Q is positive semi-definite, i.e. $Q \succeq 0$,
2. all eigenvalues of Q are non-negative, i.e. $\lambda_i(Q) \geq 0$ for all $i \in [1, n]$,
3. all principal minors of Q , i.e. the determinant of a submatrix obtained from Q when the same set of rows and columns are stricken out, are non-negative,
4. Q can be written as $Q = AA^T$ for some matrix $A \in \mathbb{R}^{n \times r}$ and r is the rank of Q .

Theorem 4.10: Symmetric positive definite matrices

A symmetric is positive definite if $B \succ 0$, i.e.,

$$\forall z \in \mathbb{R}^n \setminus \{0\} : z^T B z > 0,$$

and the set of symmetric positive definite matrices is denoted by \mathbb{S}_{++}^n .

Theorem 4.11: Cholesky factorization

If $A \in \mathbb{R}^{n \times n}$ is symmetric positive definite, then there exists a unique lower triangular matrix $L \in \mathbb{R}^{n \times n}$ with positive diagonal entries such that

$$A = LL^T.$$

Theorem 4.12: Matrix norms

For a matrix $A \in \mathbb{R}^{m \times n}$ and two vector norms $\|\cdot\|_p$ and $\|\cdot\|_q$, the induced matrix norm is defined as

$$\|A\|_{p,q} = \max_x \{\|Ax\|_q \mid \|x\|_p \leq 1\}.$$

When $p = q$, we simply write $\|A\|_p$.

Example 4.2: Spectral norm

The ℓ_2 -induced norm or spectral norm of a matrix $A \in \mathbb{R}^{m \times n}$ is defined as

$$\|A\|_2 = \sqrt{\lambda_{\max}(A^T A)} = \sigma_{\max}(A),$$

where $\lambda_{\max}(A^T A)$ denotes the largest eigenvalue of $A^T A$ and $\sigma_{\max}(A)$ is the largest singular value of A .

Example 4.3: Frobenius norm

The Frobenius norm of a matrix $A \in \mathbb{R}^{m \times n}$ is defined as

$$\|A\|_F = \sqrt{\sum_{i=1}^m \sum_{j=1}^n a_{ij}^2} = \sqrt{\text{tr}(A^T A)}.$$

5 Sequences

Lemma 5.1: Convergence

If $\{x_k\}$ is a non-increasing and bounded below or non-decreasing and bounded above sequence, it converges to a finite real number.

Theorem 5.1: $O(\cdot)$

For a function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, we write

$$f(x) = O(g(x))$$

if and only if there exists a constant $C > 0$ and a neighborhood \mathcal{N} of 0 such that

$$\forall x \in \mathcal{N} : \|f(x)\| \leq Cg(x),$$

i.e. “ f shrinks as fast as g ”.

Theorem 5.2: $o(\cdot)$

For a function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, we write

$$f(x) = o(g(x))$$

if and only if there exists a neighborhood \mathcal{N} of 0 and a function $c : \mathcal{N} \rightarrow \mathbb{R}$ with $\lim_{x \rightarrow 0} c(x) = 0$ such that

$$\forall x \in \mathcal{N} : \|f(x)\| \leq c(x)g(x),$$

i.e. “ f shrinks faster than g ”.

6 Differential Calculus

Theorem 6.1: Lipschitz continuity

A mapping $F : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is Lipschitz continuous with Lipschitz constant L if

$$\|F(x) - F(y)\| \leq L\|x - y\| \quad \forall x, y \in \mathbb{R}^n.$$

Theorem 6.2: Linear mapping

A mapping $F : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is linear if

$$\forall x, y \in \mathbb{R}^n, \forall \lambda_1, \lambda_2 \in \mathbb{R} : F(\lambda_1 x + \lambda_2 y) = \lambda_1 F(x) + \lambda_2 F(y).$$

Theorem 6.3: Affine mapping

A mapping $F : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is affine if it is the sum of a linear mapping and a constant vector, i.e.

$$\exists A \in \mathbb{R}^{m \times n}, \exists b \in \mathbb{R}^m : F(x) = Ax + b.$$

Theorem 6.4: Quadratic function

A function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is quadratic if it can be written as

$$f(x) = \frac{1}{2}x^T Qx + q^T x + c$$

for some matrix $Q \in \mathbb{R}^{n \times n}$, vector $q \in \mathbb{R}^n$, and scalar $c \in \mathbb{R}$.

Theorem 6.5: First-order Taylor expansion

Suppose that $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is continuously differentiable at $x \in \mathbb{R}^n$. Then for all $y \in \mathbb{R}^n$, it holds that

$$f(y) = f(x) + \nabla f(x)^T(y - x) + o(\|y - x\|).$$

Lemma 6.1: Mean-value theorem

Suppose that $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is differentiable. Then for every $x, y \in \mathbb{R}^n$ there exists a $\tau \in (0, 1)$ such that

$$f(y) = f(x) + \nabla f(x + \tau(y - x))^T(y - x),$$

Moreover, if f is continuously differentiable, then

$$f(y) = f(x) + \int_0^1 \nabla f(x + t(y - x))^T(y - x) dt$$

Theorem 6.6: Hessian

The Hessian of a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is the matrix of second partial derivatives, i.e.

$$\nabla^2 f(x) = \left[\frac{\partial^2 f(x)}{\partial x_i \partial x_j} \right]_{i,j=1}^n$$

Theorem 6.7: Second-order Taylor expansion

Suppose that $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is twice continuously differentiable at $x \in \mathbb{R}^n$. Then for all $y \in \mathbb{R}^n$, it holds that

$$f(y) = f(x) + \nabla f(x)^T(y - x) + \frac{1}{2}(y - x)^T \nabla^2 f(x)(y - x) + o(\|y - x\|^2).$$

Lemma 6.2: Taylor Rest Term theorem

Suppose that $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is continuously differentiable. Then for every $x, y \in \mathbb{R}^n$ there exists a $\theta \in [0, 1]$ such that

$$f(y) = f(x) + \nabla f(x)^T(y - x) + \frac{1}{2}(y - x)^T \nabla^2 f(x + \theta(y - x))(y - x).$$

Theorem 6.8: Jacobian

The Jacobian of a mapping $F : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is the matrix of transposed gradients, i.e.

$$J_F(x) = \begin{bmatrix} \nabla f_1(x)^\top \\ \nabla f_2(x)^\top \\ \vdots \\ \nabla f_m(x)^\top \end{bmatrix} \in \mathbb{R}^{m \times n}$$

Application 6.1: Mean-value theorem for vector-valued functions

Suppose that $F : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is continuously differentiable on \mathbb{R}^n . Then for every $x, y \in \mathbb{R}^n$ the following holds:

$$F(y) = F(x) + \int_0^1 J_F(x + t(y - x))(y - x) dt.$$

Theorem 6.9: Implicit function theorem

Let $F : \mathbb{R}^{n+m} \rightarrow \mathbb{R}^m$ be a continuously differentiable mapping of $x \in \mathbb{R}^n$ and $p \in \mathbb{R}^m$. If $x^* \in \mathbb{R}^n, p^* \in \mathbb{R}^m$ are such that

1. $F(x^*, p^*) = 0$,
2. the partial Jacobian $J_{F_x}(x^*, p^*)$ is non-singular,

then there exist open sets $S_{x^*} \subset \mathbb{R}^n, S_{p^*} \subset \mathbb{R}^m$ and a continuously differentiable function $g : S_{p^*} \rightarrow S_{x^*}$ such that

$$x^* = g(p^*) \quad \text{and} \quad F(g(p), p) = 0 \quad \forall p \in S_{p^*},$$

and

$$J_g(p^*) = -(J_{F_x}(g(x^*), p^*))^{-1} J_{F_p}(g(x^*), p^*).$$

Fundamental Concepts

7 Fundamental Concepts of Optimization

Theorem 7.1: Optimization problem in standard form

$$\begin{aligned} & \underset{x \in \mathbb{R}^n}{\text{minimize}} && f(x) \\ & \text{subject to} && g(x) = 0, \\ & && h(x) \geq 0. \end{aligned}$$

Theorem 7.2: Level Set

The set

$$\{x \in \mathbb{R}^n \mid f(x) = c\}$$

is the level set of f for the value c , i.e. the set of all points that map to the same value c .

Theorem 7.3: Feasible set

The set

$$\{x \in \mathbb{R}^n \mid g(x) = 0 \wedge h(x) \geq 0\}$$

is the feasible set Ω , i.e. the set of all points that satisfy the constraints.

Theorem 7.4: Global minimizer

The point x^* is a global minimizer if and only if

$$x^* \in \Omega, \forall x \in \Omega : f(x^*) \leq f(x).$$

Theorem 7.5: Strict global minimizer

The point x^* is a strict global minimizer if and only if

$$x^* \in \Omega, \forall x \in \Omega \setminus \{x^*\} : f(x^*) < f(x).$$

Theorem 7.6: Strict local minimizer

The point x^* is a strict local minimizer if and only if $x^* \in \Omega$ and there exists a neighborhood \mathcal{N} of x^* such that

$$\forall x \in (\Omega \cap \mathcal{N}) \setminus \{x^*\} : f(x^*) < f(x).$$

Theorem 7.7: Weierstrass

If $\Omega \in \mathbb{R}^n$ is compact, i.e. limited and closed, and $f : \Omega \rightarrow \mathbb{R}$ is continuous, then there exists a global minimizer (a solution) of the optimization problem

$$\begin{array}{ll} \text{minimize} & f(x) \\ & x \in \mathbb{R}^n \\ \text{subject to} & x \in \Omega. \end{array}$$

Proof 7.1: Weierstrass

Regard the graph of f , $G = \{(x, f(x)) \in \mathbb{R}^n \times \mathbb{R} \mid x \in \Omega\}$. G is a compact set, and so is the projection of G onto its last coordinate, the set $G = \{s \in \mathbb{R} \mid \exists x \text{ such that } (x, s) \in G\}$, which is a compact interval $[f_{\min}, f_{\max}] \subset \mathbb{R}$. By construction, there must be at least one x^* such that $(x^*, f(x^*)) \in G$. □

8 Types of Optimization Problems

Theorem 8.1: Convex set

A set $\Omega \subset \mathbb{R}^n$ is convex if

$$\forall x, y \in \Omega, \forall \lambda \in [0, 1] : x + \lambda(y - x) \in \Omega.$$

or if “all connecting lines lie inside the set”.

Theorem 8.2: Convex function

A function $f : \Omega \rightarrow \mathbb{R}$ is convex if

$$\forall x, y \in \Omega, \forall \lambda \in [0, 1] : f(x + \lambda(y - x)) \leq f(x) + \lambda(f(y) - f(x)).$$

or if “all secants (i.e. a line segment between two points on the graph) are above graph”. This definition is equivalent to saying that the Epigraph of f , i.e. the set $\{(x, s) \in \mathbb{R}^n \times \mathbb{R} \mid x \in \Omega, s \geq f(x)\}$, is a convex set.

Note: a concave function is the same but then with \geq instead of \leq .

Property 8.1: Convex function

If $f : D \rightarrow \mathbb{R}$ and $\Omega_f = \{(x, y) \mid x \in D, y \geq f(x)\}$ then the following holds:

$$f \text{ is convex} \Leftrightarrow \Omega_f \text{ is convex.}$$

Property 8.2: Globality of local minima of convex function


For a convex function, a local minimum is also a global one.

Proof 8.1: Globality of local minima of convex function

First we choose, using local optimality, a neighborhood \mathcal{N} of x^* such that for all $\tilde{x} \in \Omega \cap \mathcal{N}$ holds $f(\tilde{x}) \geq f(x^*)$. Second, we regard the connecting line between x^* and y . This line is completely contained in Ω due to the convexity of Ω . Now we choose a point \tilde{x} on this line such that it is in the neighborhood, but not equal to x^* , i.e. $\tilde{x} = x^* + \lambda(y - x^*)$ for some $\lambda \in (0, 1)$, and $\tilde{x} \in \Omega \cap \mathcal{N}$. Due to local optimality, we have $f(\tilde{x}) \geq f(x^*)$, and due to convexity we have

$$f(\tilde{x}) = f(x^* + \lambda(y - x^*)) \leq f(x^*) + \lambda(f(y) - f(x^*)).$$

It follows that $\lambda(f(y) - f(x^*)) \geq 0$, and since $\lambda \in (0, 1)$, we have $f(y) \geq f(x^*)$.



□

Example 8.1: Nonlinear Programming (NLP)

Nonlinear Programming Problems are problems of the following form:

$$\begin{aligned} & \underset{x \in \mathbb{R}^n}{\text{minimize}} && f(x) \\ & \text{subject to} && g(x) = 0, \\ & && h(x) \geq 0, \end{aligned}$$

where $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $g : \mathbb{R}^n \rightarrow \mathbb{R}^p$, and $h : \mathbb{R}^n \rightarrow \mathbb{R}^q$, are assumed to be continuously differentiable at least once.

Example 8.2: Linear Programming (LP)

When the functions f, g, h are affine (i.e., they can be expressed in the form $f(x) = a^T x + b$ for some vector a and scalar b) in the general formulation (see Theorem 8.1), the general NLP gets something easier to solve, namely a Linear Program (LP), which can be written as follows:

$$\begin{aligned} & \underset{x \in \mathbb{R}^n}{\text{minimize}} && c^T x \\ & \text{subject to} && Ax - b = 0, \\ & && Cx - d \geq 0, \end{aligned}$$

where $c \in \mathbb{R}^n$, $A \in \mathbb{R}^{p \times n}$, $b \in \mathbb{R}^p$, $C \in \mathbb{R}^{q \times n}$, and $d \in \mathbb{R}^q$.

Example 8.3: Quadratic Programming (QP)

When the functions g, h are affine (as for an LP in Theorem 8.2), and the objective function f is a linear-quadratic function, the NLP becomes a Quadratic Program (QP), which can be written as follows:

$$\begin{aligned} & \underset{x \in \mathbb{R}^n}{\text{minimize}} && c^T x + \frac{1}{2} x^T B x \\ & \text{subject to} && Ax - b = 0, \\ & && Cx - d \geq 0, \end{aligned}$$

where, in addition to the LP parameters, $B \in \mathbb{R}^{n \times n}$ is the Hessian matrix. Specifically, for the objective function $f(x) = c^T x + \frac{1}{2} x^T B x$, the Hessian is given by:

$$\nabla^2 f(x) = B.$$

Theorem 8.3: Convex QP

If the Hessian matrix B is positive semi-definite (i.e. if $\forall z \in \mathbb{R}^n \setminus \{0\} : z^T B z \geq 0$) we call the QP a convex QP. Convex QPs are tremendously easier to solve globally than “non-convex QPs” (i.e., where the Hessian B is not positive semi-definite), which might have different local minima (i.e. have a non-convex solution set).

Theorem 8.4: Strictly convex QP


If the Hessian matrix B is positive definite (i.e. if $\forall z \in \mathbb{R}^n : z^T B z > 0$) we call the QP a strictly convex QP. Strictly convex QPs are a subset of convex QPs (see Theorem 8.3).

Example 8.4: Mixed-Integer Programming (MIP)

A mixed-integer programming problem or mixed-integer program (MIP) is a problem with both real and integer decision variables. A MIP can be formulated as follows:

$$\begin{aligned} & \underset{\substack{x \in \mathbb{R}^n \\ z \in \mathbb{Z}^m}}{\text{minimize}} && f(x, z) \\ & \text{subject to} && g(x, z) = 0, \\ & && h(x, z) \geq 0, \end{aligned}$$

where $x \in \mathbb{R}^n$ are the continuous variables, and $z \in \mathbb{Z}^m$ are the integer variables.




9 Convex Optimization

Theorem 9.1: Convexity for C^1 functions

Assume that $f : \Omega \rightarrow \mathbb{R}$ is continuously differentiable and Ω is convex. Then holds that f is convex if and only if

$$\forall x, y \in \Omega : f(y) \geq f(x) + \nabla f(x)^T(y - x)$$

i.e. tangents lie below the graph.



Proof 9.1: Convexity for C^1 functions

“ \Rightarrow ”: Due to convexity of f holds for given $x, y \in \Omega$ and for any $\lambda \in [0, 1]$ that

$$f(x + \lambda(y - x)) - f(x) \leq \lambda(f(y) - f(x)),$$

and therefore that

$$\frac{f(x + \lambda(y - x)) - f(x)}{\lambda} \leq f(y) - f(x).$$

Furthermore, we can deduce that

$$\lim_{\lambda \rightarrow 0} \frac{f(x + \lambda(y - x)) - f(x)}{\lambda} = \nabla f(x)^T(y - x) \leq f(y) - f(x),$$

which proves the statement.

“ \Leftarrow ”: To prove that for $z = x + \lambda(y - x) = (1 - \lambda)x + \lambda y$ holds that $f(z) \leq (1 - \lambda)f(x) + \lambda f(y)$, we can use the equation from Theorem 9.1 twice to get

$$f(x) \geq f(z) + \nabla f(z)^T(x - z) \quad \text{and} \quad f(y) \geq f(z) + \nabla f(z)^T(y - z),$$

which yield, when weighted with $(1 - \lambda)$ and λ respectively, that

$$(1 - \lambda)f(x) + \lambda f(y) \geq f(z) + \nabla f(z)^T \underbrace{[(1 - \lambda)(x - z) + \lambda(y - z)]}_{=0}$$

□

Theorem 9.2: Convexity for C^2 functions

Assume that $f : \Omega \rightarrow \mathbb{R}$ is twice continuously differentiable and Ω is convex and open. Then holds that f is convex if and only if

$$\forall x \in \Omega : \nabla^2 f(x) \succeq 0,$$

i.e. the Hessian of f is positive semi-definite.

Proof 9.2: Convexity for C^2 functions

“ \Rightarrow ”: Recall that a second order Taylor expansion of y at x in an arbitrary direction p is given by the following:

$$f(x + tp) = f(x) + t\nabla f(x)^T p + \frac{t^2}{2} p^T \nabla^2 f(x) p + o(t^2 \|p\|).$$

From this we obtain that

$$p^T \nabla^2 f(x) p = \lim_{t \rightarrow 0} \frac{2}{t^2} \left(\underbrace{f(x + tp) - f(x) - t\nabla f(x)^T p}_{(9.1): \geq 0} \right) \geq 0.$$

“ \Leftarrow ”: Conversely, to prove the other direction, we use Theorem 6.2 with some arbitrary $\theta \in [0, 1]$:

$$f(y) = f(x) + \nabla f(x)^T (y - x) + \underbrace{\frac{t^2}{2} (y - x)^T \nabla^2 f(x + \theta(y - x)) (y - x)}_{(9.2): \geq 0}.$$

□

Property 9.1: Convexity perserving operations on convex functions

The following operations preserve the convexity of a function:

1. Non-negative weighted sum: Suppose that $\forall i \in [1, m] : f_i : \mathbb{R}^n \rightarrow \mathbb{R}$ are convex functions and $\forall i \in [1, m] : \lambda_i \geq 0$. Then the following function is convex:

$$f(x) = \sum_{i=1}^m \lambda_i f_i(x)$$

2. Affine input transformation: If $f : \Omega \rightarrow \mathbb{R}$ is convex, then

$$A \in \mathbb{R}^{n \times m} : \tilde{f}(x) = f(Ax + b)$$

is convex on the domain $\tilde{\Omega} = \{x \mid Ax + b \in \Omega\}$.

3. Concatenation with monotone convex function: If $f : \Omega \rightarrow \mathbb{R}$ is convex and $g : \mathbb{R} \rightarrow \mathbb{R}$ is convex and monotonely increasing, then the composition $g \circ f$ is convex.
4. Pointwise supremum: The supremum over a set of convex functions $f_i(x)$, $i \in I$, where I can be an infinite set, i.e.,

$$f(x) = \sup_{i \in I} f_i(x)$$

is convex.

5. Composition: Let $h : \mathbb{R}^m \rightarrow \mathbb{R}$ and $g : \mathbb{R}^n \rightarrow \mathbb{R}^m$ with $g = (g_1, \dots, g_m)$. Then $f(x) = (h \circ g)(x) = h(g(x))$ is convex if any of the followings holds:

- (a) h is convex and non-decreasing in each argument and $\forall i \in [1, m] : g_i$ is convex.
- (b) h is convex and non-decreasing in each argument and $\forall i \in [1, m] : g_i$ is concave, i.e. $-g_i$ is convex.

Proof 9.3: Convexity perserving operations on convex functions

1. ...
2. ...
3. Recall that g is a convex and monotonely increasing function, then:

$$\nabla^2(g \circ f) = \underbrace{g''(f(x))}_{\geq 0} \underbrace{\nabla f(x) \nabla f(x)^T}_{\succeq 0} + \underbrace{g'(f(x))}_{\geq 0} \underbrace{\nabla^2 f(x)}_{\succeq 0} \succeq 0,$$

i.e. $g \circ f$ is convex, since the Hessian is positive semi-definite.

4. Epigraph of f is the intersection of the epigraphs of f_i , which are convex.
5. Recall that g_i is convex and that h is convex and non-decreasing in each argument. Then:

$$\begin{aligned} f(\lambda x + (1 - \lambda)y) &= h(g(\lambda x + (1 - \lambda)y)) \\ &\leq h(\lambda g(x) + (1 - \lambda)g(y)) \end{aligned} \tag{1}$$

$$\begin{aligned} &\leq \lambda h(g(x)) + (1 - \lambda)h(g(y)) \\ &= \lambda f(x) + (1 - \lambda)f(y), \end{aligned} \tag{2}$$


where we used the convexity of g_i and the fact that h is non-decreasing to obtain inequality (1), and the convexity of h to obtain inequality (2). □

Theorem 9.3: Convexity of Sublevel subsets

If $f : \Omega \rightarrow \mathbb{R}$ is a convex function, then all its level sets

$$\text{lev}_{\leq \gamma} f = \{x \in \Omega \mid f(x) \leq \gamma\}$$

are convex.



Proof 9.4: Convexity of Sublevel subsets

If $f(x) \leq c$ and $f(y) \leq c$, then for any $\lambda \in [0, 1]$ holds also

$$f((1 - \lambda)x + \lambda y) \leq (1 - \lambda)f(x) + \lambda f(y) \leq \underbrace{(1 - \lambda)c + \lambda c}_{=c}.$$

□

Property 9.2: Convexity perserving operations on convex sets

The following operations preserve the convexity of a set:

1. The intersection of finitely or infinitely many convex sets is convex.
2. Affine image: if Ω is convex, then for $A \in \mathbb{R}^{m \times n}$, $b \in \mathbb{R}^m$ also the set $A\Omega + b = \{y \in \mathbb{R}^m \mid \exists x \in \Omega, y = Ax + b\}$ is convex.

$\Omega : y = Ax + b\}$ is convex.

3. Affine pre-image: if Ω is convex, then for $A \in \mathbb{R}^{n \times m}$, $b \in \mathbb{R}^n$ also the set $\{z \in \mathbb{R}^m \mid Az + b \in \Omega\}$ is convex.

Example 9.1: Convex Feasible set

If $\forall i \in [1, m] : f_i : \mathbb{R}^n \rightarrow \mathbb{R}$ are convex functions, then the set

$$\Omega = \{x \in \mathbb{R}^n \mid f_i(x) \leq 0, i \in [1, m]\}$$

is a convex set, because it is the intersection of sublevel sets Ω_i of convex functions f_i , i.e.

$$\Omega = \bigcap_{i=1}^m \{x \in \mathbb{R}^n \mid f_i(x) \leq 0\}.$$

Theorem 9.4: Optimality condition for convex problems

Regard a convex optimization problem with continuously differentiable objective function f . A point $x^* \in \Omega$ is a global optimizer if and only if

$$\forall y \in \Omega : \nabla f(x^*)^T (y - x^*) \geq 0.$$

Proof 9.5: Optimality condition for convex problems

“ \Rightarrow ”: Assume for the sake of contradiction that $\exists y \in \Omega : \nabla f(x^*)^T (y - x^*) < 0$, then we could regard a Taylor expansion

$$f(x^* + \lambda(y - x^*)) = f(x^*) + \underbrace{\lambda \nabla f(x^*)^T (y - x^*)}_{<0} + \underbrace{o(\lambda)}_{\rightarrow 0}.$$

This implies that for sufficiently small positive λ , we have

$$f(x^* + \lambda(y - x^*)) < f(x^*),$$

which contradicts the optimality of x^* .

“ \Leftarrow ”: Due to the C^1 characterization of convexity of f in Theorem 9.1, we have for any feasible $y \in \Omega$:

$$f(y) \geq f(x^*) + \underbrace{\nabla f(x^*)^T (y - x^*)}_{\geq 0} \geq f(x^*),$$

which implies that x^* is a global optimizer.

□

Theorem 9.5: Sufficient Condition for Convex NLP

For a nonlinear optimization problem (NLP) in standard form

$$\begin{aligned} & \underset{x \in \mathbb{R}^n}{\text{minimize}} && f(x) \\ & \text{subject to} && g_i(x) = 0, \ i \in [1, m], \\ & && h_i(x) \geq 0, \ j \in [1, p], \end{aligned}$$

the following conditions are necessary for convexity:

- the objective function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ must be convex,
- the constraint set

$$X = \{x \in \mathbb{R}^n \mid g(x) = 0, \ h(x) \geq 0\}$$

must be convex. Since we know that the intersection of convex sets is convex, we can write X as the intersection of the sets G and H :

$$\begin{aligned} X &= \{x \in \mathbb{R}^n \mid g(x) = 0, \ h(x) \geq 0\} \\ &= \{x \in \mathbb{R}^n \mid g(x) = 0\} \cap \{x \in \mathbb{R}^n \mid h(x) \geq 0\} \\ &= \left(\bigcap_{i=1}^m \{x \in \mathbb{R}^n \mid g_i(x) = 0\} \right) \cap \left(\bigcap_{i=1}^m \{x \in \mathbb{R}^n \mid h_i(x) \geq 0\} \right) \\ &= G \cap H \end{aligned}$$

Now we must consider the requirements for the sets G and H to be convex:

- Suppose $-H_i = \{x \in \mathbb{R}^n \mid h_i(x) \leq 0\}$, the zero sublevel set of the function h_i . If h_i is convex, the set $-H_i$ is convex, as seen in Theorem 9.3. Now, consider the case where h_i is concave. Since h_i being concave means that $-h_i$ is convex, the set $H_i = \{x \in \mathbb{R}^n \mid h_i(x) \geq 0\}$ is a sublevel set of the convex function $-h_i$, and is therefore convex. Thus, the set H_i is convex when h_i is concave.
- On the other hand, G is the level set of g_i . Therefore it is certainly a convex set whenever g_i is a affine function

$$\forall i \in [1, m] : g_i(x) = a_i^T x + b_i.$$

Example 9.2: Halfspace

A halfspace

$$H_{\leq} = \{x \in \mathbb{R}^n \mid a^T x \leq b\}$$

is a convex set, as the zero sublevel set of the affine (and therefore convex) function $f(x) = a^T x - b$, cf. Theorem 9.3.

Note: The opposite is not true; a function that has all its level sets convex is not necessarily convex.

Example 9.3: Polyhedral Set

A polyhedral set $C \subset \mathbb{R}^n$ is defined as the intersection of a finite number of halfspaces, i.e.,

$$C = \bigcap_{i=1, \dots, m} \{x \in \mathbb{R}^n \mid a_i^T x \leq b_i\}.$$

Since the intersection of convex sets is convex, C is a convex set if all the halfspaces are convex.

Note: A polyhedral set might contain equalities, i.e.,

$$C = \{x \in \mathbb{R}^n \mid Ax \leq b, Cx = d\}.$$

which can be written with just inequalities as such:

$$C = \{x \in \mathbb{R}^n \mid Ax \leq b, Cx \leq d, -Cx \leq -d\}.$$

Example 9.4: Ellipsoid

If $P \in \mathbb{R}^{n \times n}$ is symmetric positive definite, then the ellipsoid

$$C = \{x \in \mathbb{R}^n \mid (x - x_c)^T P^{-1} (x - x_c) \leq 1\}$$

is convex, as the 1-sublevel set of the convex function $f(x) = (x - x_c)^T P^{-1} (x - x_c) - 1$, where x_c is the center of the ellipsoid and P is the shape matrix. The latter determines how far the ellipsoid extends in each direction from x_c ; the lengths of the semi-axis are given by $\sqrt{\lambda_i}$, where λ_i are the eigenvalues of P , while the eigenvectors of P determine the orientation of the ellipsoid. When $P = r^2 I$, the ellipsoid is a ball with radius r around x_c , i.e.,


$$C = \{x \in \mathbb{R}^n \mid \|x - x_c\|_2 \leq r\},$$

or, in general, the p -norm ball


$$C = \{x \in \mathbb{R}^n \mid \|x - x_c\|_p \leq r\}.$$

where the value of p determines the shape of the ball, i.e.,


- $p = 1$:



- $p = 2$:



- $p = \infty$:



Note: All images above are for $r = 1$.

Example 9.5: Convex cones

A set C is said to be cone if


$$\forall x : x \in C \Rightarrow \forall \lambda \geq 0 : \lambda x \in C.$$

Moreover, C is a convex cone if and only if it is closed under addition and multiplication with non-negative scalars, i.e.,

$$\forall x, y \in C, \forall \lambda, \mu \geq 0 : \lambda x + \mu y \in C.$$

The following are examples of convex cones:

- Non-negative orthant: The set $C = \{x \in \mathbb{R}^n \mid x \geq 0\}$ is a convex cone.
- Norm cones: The set $C = \{(x, t) \in \mathbb{R}^n \times \mathbb{R} \mid \|x\|_p \leq t\}$ is a convex cone. For instance, the Lorenz cone or ice-cream cone, when $p = 2$, is shown on the right.
- Positive semi-definite cone: The set \mathbb{S}_+^n is a convex cone.



10 The Lagrangian Function and Duality

Theorem 10.1: Primal Optimization Problem

We will denote the globally optimal value of the objective function subject to the constraints as the primal optimal value p^* , i.e.,

$$p^* = \underbrace{\left(\min_{x \in \mathbb{R}^n} f(x) \quad \text{s.t.} \quad g(x) = 0 \wedge h(x) \geq 0 \right)}_{\text{primal optimization problem}}.$$

Theorem 10.2: Lagrangian Function and Lagrange Multipliers

We define the Lagrangian function to be

$$\begin{aligned} \mathcal{L}(x, \lambda, \mu) &= f(x) + \sum_{i=1}^l \lambda_i g_i(x) + \sum_{i=1}^m \mu_i h_i(x) \\ &= f(x) + \lambda^T g(x) + \mu^T h(x). \end{aligned}$$

where $\lambda \in \mathbb{R}^l$ and $\mu \geq 0 \in \mathbb{R}^m$ are the Lagrange multipliers or dual variables.

Lemma 10.1: Lower Bound Property of Lagrangian

If \tilde{x} is a feasible point of (10.1) and $\mu \geq 0$, then

$$\mathcal{L}(\tilde{x}, \lambda, \mu) \leq f(\tilde{x}).$$

Proof 10.1: Lower Bound Property of Lagrangian

Since \tilde{x} is feasible, we have $g(\tilde{x}) = 0$ and $h(\tilde{x}) \geq 0$. Therefore, with $\mu \geq 0$, we have:

$$\mathcal{L}(\tilde{x}, \lambda, \mu) = f(\tilde{x}) + \lambda^T \underbrace{g(\tilde{x})}_{=0} + \underbrace{\mu^T}_{\geq 0} \underbrace{h(\tilde{x})}_{\geq 0} \leq f(\tilde{x}).$$

□

Theorem 10.3: Lagrange Dual Function

The Lagrange dual function is defined as the unconstrained infimum of the Lagrangian function over x , for fixed multipliers λ and μ :

$$q(\lambda, \mu) = \inf_{x \in \mathbb{R}^n} \mathcal{L}(x, \lambda, \mu).$$

Lemma 10.2: Lower Bound property of Lagrange Dual

If $\mu \geq 0$, then the dual function $q(\lambda, \mu)$ is a lower bound on the primal optimal value p^* , i.e.,

$$q(\lambda, \mu) \leq p^*.$$

Proof 10.2: Lower Bound property of Lagrange Dual

Since the Lagrange function is bounded from below by Lemma 10.1, we have

$$q(\lambda, \mu) = \inf_{x \in \mathbb{R}^n} \mathcal{L}(x, \lambda, \mu) \leq f(\tilde{x}) \quad \text{for any feasible } \tilde{x}.$$

Naturally, this inequality holds in particular for the global minimizer x^* , which yields:

$$q(\lambda, \mu) \leq f(x^*) = p^*.$$


□

Theorem 10.4: Concavity of Lagrange Dual

The function $q : \mathbb{R}^p \times \mathbb{R}^q \rightarrow \mathbb{R}$ is concave, even if the original NLP was not convex.

Proof 10.3: Concavity of Lagrange Dual

We will show that $-q$ is convex. The Lagrangian \mathcal{L} is an affine function in the multipliers λ and μ , which in particular implies that $-\mathcal{L}$ is convex in λ and μ . Thus, the function $-q(\lambda, \mu) = \sup_x -\mathcal{L}(x, \lambda, \mu)$ is the supremum of convex functions in λ and μ that are indexed by x , and therefore convex.



Theorem 10.5: Dual Problem

The dual problem is defined as the convex maximization problem, i.e.,

$$d^* = \left(\max_{\lambda \in R^p, \mu \in R^q} q(\lambda, \mu) \quad \text{s.t.} \quad \mu \geq 0 \right).$$

Theorem 10.6: Weak Duality

Consider a primal-dual pair. Then, the following inequality holds:

$$d^* \leq p^*.$$

Theorem 10.7: Strong Duality

If the primal optimization problem is convex and the Slater condition (see Theorem 10.8) holds, then strong duality holds, i.e.,

$$d^* = p^*.$$

Theorem 10.8: Slater Condition

If there exist one feasible point \tilde{x} such that all non-linear inequalities are strictly satisfied of a primal convex optimization problem hold, then the Slater condition is satisfied. More explicitly, for a convex problem we must have affine equality constraints, $g(x) = Ax + b$, and the inequality constraint functions can be either affine or concave functions, thus without loss of generality assume that the first $q_1 \leq q$ inequalities are affine and the remaining ones concave. Then the Slater condition holds if and only if there exists an \tilde{x} such that

$$\begin{aligned} A\tilde{x} + b &= 0, \\ h_i(\tilde{x}) &\geq 0, \quad \text{for } i = 1, \dots, q_1, \\ h_i(\tilde{x}) &> 0, \quad \text{for } i = q_1 + 1, \dots, q. \end{aligned}$$

Note: This is trivially satisfied for LP and QP problems.

Theorem 10.9: KKT

Consider the convex optimization problem with equality and inequality constraints, i.e.,

$$\begin{aligned} \min_x \quad & f(x) \\ \text{s.t.} \quad & g(x) = 0, \\ & h(x) \geq 0, \end{aligned}$$

and assume that f is convex, g is affine, and h is concave. Also assume that Slater's condition holds and all functions are differentiable. Then, the following statements are equivalent:

1. x^* is a primal optimal and (λ^*, μ^*) are dual optimal.
2. (x^*, λ^*, μ^*) satisfy the Karush-Kuhn-Tucker (KKT) conditions:

1. Stationarity:

$$\nabla f(x) - \nabla g(x)\lambda - \nabla h(x)\mu = 0$$

2. Primal feasibility:

$$g(x) = 0$$

3. Dual feasibility:

$$h(x) \geq 0$$

4. Non-negativity of Lagrange multipliers:

$$\mu \geq 0$$

5. Complementary slackness:

$$\mu^T h(x) = 0$$

Proof 10.4: KKT

“ \Rightarrow ”: Because of the assumptions, there is no duality gap. Thus, there is strong duality and we have:

$$\begin{aligned} p^* = d^* &= q(\lambda^*, \mu^*) \\ &= \inf_x \mathcal{L}(x, \lambda^*, \mu^*) \\ &= \inf_x f(x) + \lambda^{*T} g(x) + \mu^{*T} h(x) \\ &\leq f(x^*) + \lambda^{*T} g(x^*) + \mu^{*T} h(x^*) \\ &= p^* - \underbrace{\mu^{*T} h(x^*)}_{\geq 0}. \end{aligned}$$

This implies that $\mu^{*T} h(x^*) = 0$ and $\nabla f(x) - \nabla g(x)\lambda^* - \nabla h(x)\mu^* = 0$. The remaining KKT conditions follow trivially.

“ \Leftarrow ”: Since the assumptions still hold and the KKT condition of stationarity is satisfied for (x^*, λ^*, μ^*) , we conclude that x^* is a global minimizer of the convex function $\mathcal{L}(x, \lambda^*, \mu^*)$. Therefore we have:

$$\begin{aligned} q(\lambda^*, \mu^*) &= \mathcal{L}(x^*, \lambda^*, \mu^*) \\ &= f(x^*) - \underbrace{\lambda^* g(x^*)}_{=0} - \underbrace{\mu^* h(x^*)}_{=0} \\ &= f(x^*), \end{aligned}$$

with the last line following from the other KKT conditions. We conclude:

$$p^* \leq f(x^*) = q(\lambda^*, \mu^*) \leq d^*.$$

The assumptions imply strong duality, and thus that $p^* = d^*$. Since we always have $d^* \leq p^*$, we conclude that the inequalities are in fact equalities. Consequently, x^* is primal optimal and (λ^*, μ^*) are dual optimal.

□

Example 10.1: Dual Decomposition

Unconstrained Optimization Algorithms

11 Optimality Conditions

Theorem 11.1: Unconstrained optimization Problems

We define the unconstrained optimization problem as

$$\min_{x \in D} f(x),$$

where we regard objective function $f : D \rightarrow \mathbb{R}$ that are defined on some open domain $D \subseteq \mathbb{R}^n$.

Theorem 11.2: Stationary Point

A point \tilde{x} is called a stationary point of f if and only if

$$\nabla f(\tilde{x}) = 0.$$

Theorem 11.3: Descent Direction

A vector $p \in \mathbb{R}^n$ is called a descent direction at x if

$$\nabla f(x)^T p < 0.$$

Theorem 11.4: First Order Necessary Conditions (FONC)

If $x^* \in D$ is a local minimizer of $f : D \rightarrow \mathbb{R}$ and $f \in C^1$ then

$$\nabla f(x^*) = 0.$$

Proof 11.1: First Order Necessary Conditions (FONC)

Let us assume for contradiction that $\nabla f(x^*) \neq 0$. Then $p = -\nabla f(x^*)$ would be a descent direction in which the objective could be improved, as follows:

As D is open and $f \in C^1$, we could find a $t > 0$ that is small enough so that for all $\tau \in [0, t]$ holds $x^* + \tau p \in D$ and $\nabla f(x^* + \tau p)^T p < 0$. By Taylor's theorem, we would have for some $\theta \in (0, 1)$ that

$$f(x^* + tp) = f(x^*) + \underbrace{t \nabla f(x^* + \theta tp)^T p}_{< 0} < f(x^*).$$

□

Theorem 11.5: Second Order Necessary Conditions (SONC)

If $x^* \in D$ is a local minimizer of $f : D \rightarrow \mathbb{R}$ and $f \in C^2$ then

$$\nabla^2 f(x^*) \succeq 0.$$

Proof 11.2: Second Order Necessary Conditions (SONC)

Assume, for the sake of contradiction, that $\nabla^2 f(x^*) \prec 0$. This implies the existence of a vector $p \in \mathbb{R}^n$ such that $p^T \nabla^2 f(x^*) p < 0$. In this case, the objective function could be improved in the direction of p . By choosing a sufficiently small $t > 0$, we can ensure that for all $\tau \in [0, t]$, the following holds:

$$p^T \nabla^2 f(x^* + \tau p) p < 0.$$

Applying Taylor's theorem, we would have for some $\theta \in (0, 1)$ that

$$f(x^* + tp) = f(x^*) + \underbrace{t \nabla f(x^*)^T p}_{=0} + \frac{t^2}{2} \underbrace{p^T \nabla^2 f(x^* + \theta p) p}_{<0} < f(x^*),$$

which leads to a contradiction. □

Theorem 11.6: Convex First Order Sufficient Conditions (cFOSC)

Assume that $f : D \rightarrow \mathbb{R}$ is convex and $f \in C^1$. If $x^* \in D$ is a stationary point of f , then x^* is a global minimizer of f .

Theorem 11.7: Second Order Sufficient Conditions (SOSC)

Assume that $f : D \rightarrow \mathbb{R}$ and $f \in C^2$. If $x^* \in D$ is a stationary point of f and $\nabla^2 f(x^*) \succ 0$, then x^* is a strict local minimizer of f .

Proof 11.3: Second Order Sufficient Conditions (SOSC)

We can choose a sufficiently small closed ball B around x^* so that for all $x \in B$ holds $\nabla^2 f(x) \succ 0$. Restricted to this ball, we have a convex problem, so that Theorem 11.6 together with stationarity of x^* implies that x^* is a global minimizer of f . To prove that it is strict, we look for any $x \in B \setminus x^*$ at the Taylor expansion, which yields with some $\theta \in (0, 1)$:

$$f(x) = f(x^*) + \underbrace{\nabla f(x^*)^T (x - x^*)}_{=0} + \frac{1}{2} \underbrace{(x - x^*)^T \nabla^2 f(x^* + \theta(x - x^*)) (x - x^*)}_{>0} > f(x^*).$$

□

Theorem 11.8: Stability of Parametric Solutions

Assume that $f : D \times \mathbb{R}^m \rightarrow \mathbb{R}$, and regard the minimization of $f(\cdot, \tilde{a})$ for a given fixed value of $\tilde{a} \in \mathbb{R}^m$. If $\tilde{x} \in D$ satisfies the SOSC (see Theorem 11.7), then there is a neighborhood $\mathcal{N} \subset \mathbb{R}^m$ around \tilde{a} such that the parametric minimizer function $x^*(a)$ is well-defined for all $a \in \mathcal{N}$ (i.e. there is a unique minimizer $x^*(a)$ for each $a \in \mathcal{N}$), is differentiable in \mathcal{N} , and $x^*(\tilde{a}) = \tilde{x}$. Its derivative at \tilde{a} is given by

$$\frac{\partial(x^*(\tilde{a}))}{\partial a} = -(\nabla_x^2 f(\tilde{x}, \tilde{a}))^{-1} \frac{\partial(\nabla_x f(\tilde{x}, \tilde{a}))}{\partial a}.$$

Moreover, each such $x^*(a)$ with $a \in \mathcal{N}$ satisfies again the SOSC and is thus a strict local minimizer.

Proof 11.4: Stability of Parametric Solutions

The existence of the differentiable map $x^* : \mathcal{N} \rightarrow D$ follows from the implicit function theorem applied to the stationarity condition $\nabla_x f(x^*(a), a) = 0$. The derivative of $x^*(a)$ at \tilde{a} is given by

$$\frac{d(\nabla_x f(x^*(a), a))}{da} = \underbrace{\frac{\partial(\nabla_x f(x^*(a), a))}{\partial x}}_{=\nabla_x^2 f} \frac{\partial x^*(a)}{\partial a} + \frac{\partial(\nabla_x f(x^*(a), a))}{\partial a} = 0$$

The fact that all points $x^*(a)$ satisfy the SOSC follows from the continuity of the second derivative. \square

12 Estimation and Fitting Problems

Theorem 12.1: Estimation and Fitting Problems

Estimation and fitting problems are optimization problems with a special objective, namely a least squares objective. We define the estimation problem as

$$\min_{x \in \mathbb{R}^n} \frac{1}{2} \|\eta - M(x)\|_2^2,$$

where $\eta \in \mathbb{R}^m$ is the measurement vector, $M : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is the model function, and $x \in \mathbb{R}^n$ is the parameter vector. Many models in estimation and fitting problems are linear functions of x . If M is linear, $M(x) = Jx$, then $f(x) = \frac{1}{2} \|\eta - Jx\|_2^2$ which is a convex function, as $\nabla^2 f(x) = J^T J \succeq 0$. Therefore local minimizers are found by:

$$\begin{aligned} \nabla f(x) = 0 &\Leftrightarrow J^T Jx^* - J^T \eta = 0 \\ &\Leftrightarrow x^* = \underbrace{(J^T J)^{-1} J^T}_{J^+} \eta \end{aligned}$$

Theorem 12.2: Pseudo-inverse

J^+ is called the pseudo-inverse and is a generalization of the inverse matrix. If $J^T J \succ 0$, J^+ is given

by

$$J^+ = (J^T J)^{-1} J^T.$$

So far, $(J^T J)^{-1}$ is only defined if $J^T J \succ 0$. This holds if and only if $\text{rank}(J) = n$, i.e. if the columns of J are linearly independent.

13 Newton Type Optimization

14 Globalisation Strategies

15 Calculating Derivatives

Constrained Optimization Algorithms

- 16 Optimality Conditions for Constrained Optimization
- 17 Equality Constrained Optimization Algorithms
- 18 Inequality Constrained Optimization Algorithms
- 19 Optimal Control Problems